

A Matlab implementation of L-BFGS-B

Brian Granzow

1 Introduction and motivation

In this report, we discuss a MATLAB implementation of the L-BFGS-B [1] algorithm for solving bound-constrained optimization problems of the form:

$$\begin{aligned} & \underset{x}{\text{minimize}} && f(x) \\ & \text{subject to} && l \leq x \leq u, \end{aligned} \tag{1}$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $x, l, u \in \mathbb{R}^n$. The L-BFGS-B algorithm is a quasi-Newton gradient-based optimization algorithm that utilizes limited memory BFGS approximations to the Hessian matrix $\frac{\partial^2 f}{\partial x_i \partial x_j}$, making it well-suited for optimization problems with a large number of design variables n .

The [original L-BFGS-B implementation](#) [3] and subsequent releases were implemented in Fortran. Since its original release, [C/C++](#) [4], [Matlab](#) [5], [Java](#) [6], and [Julia](#) [7] wrappers for the L-BFGS-B Fortran library have been released. Up to this point, however, an open-source pure Matlab implementation of the algorithm has been difficult to come by. This is potentially due to performance reasons, as the L-BFGS-B algorithm concerns itself mainly with large scale optimization problems, for which compiled low-level languages will almost certainly outperform Matlab.

Nevertheless, the aim of this report is to discuss a recently developed Matlab implementation of the algorithm that provides Matlab users a convenient way to solve box-constrained optimization problems with L-BFGS-B using only a single file Matlab implementation *LBFGS.m*, removing the additional step of installing a third-party library. This implementation is readily available for use at:

- <https://github.com/bgranzow/L-BFGS-B>

The remainder of this report is structured as follows. We begin by presenting a high-level overview of the L-BFGS-B algorithm, while making

appropriate observations about its Matlab implementation. Then we discuss several basic regression tests that have been implemented to ensure the L-BFGS-B Matlab implementation solves bound-constrained optimization problems appropriately.

2 The L-BFGS-B algorithm

From a high-level the L-BFGS-B algorithm iterates over quasi-Newton steps. For a given iteration k , the objective function is approximated by a quadratic model at a point x_k as:

$$m_k(x) = f_k + g_k^T(x - x_k) + \frac{1}{2}(x - x_k)^T B_k(x - x_k), \quad (2)$$

where f_k represents the objective function f evaluated at x_k , g_k denotes the objective gradient evaluated at x_k , and B_k represents a limited memory BFGS approximation to the Hessian evaluated at x_k . Using this quadratic model, the L-BFGS-B algorithm can be outlined using the following steps:

1. Check if the inf-norm of the gradient g_k projected onto the feasible design space is less than a user-specified tolerance. If it is, then return successfully.
2. Find the Cauchy point x^c , that minimizes m_k in the steepest descent direction $-g_k$ projected onto the feasible design space. Once found, the Cauchy point x^c is used to identify active design variables $A(x)$ (those that are identified as fixed at either an upper or lower bound) and free variables $F(x)$ (those that are identified as inside the feasible design space). Conceptually, this is like the algorithm ‘peeking ahead’ an iteration to determine which variables will likely be active and inactive.
3. The quadratic model m_k is then minimized for the set of free variables $F(x)$ in an unconstrained manner, and then backtracked into the feasible design space to obtain \bar{x}_k .
4. The new search direction is computed as $d_k = \bar{x}_k - x_k$ and a line-search method is used to find a step length α_k that satisfies the strong-Wolfe conditions to compute the new design variables x_{k+1} .
5. The L-BFGS Hessian approximation B_{k+1} is computed based on the new step x_{k+1} and a new iteration is started.

These steps are now discussed from a high-level, with references to their implementation.

2.1 Convergence criteria

Convergence is satisfied if the following criteria holds:

$$\|P(x_k - g_k, l, u) - x_k\|_\infty < \text{tol}, \quad (3)$$

where tol is a user-specified tolerance and

$$P(x, l, u) := \begin{cases} l_i & \text{if } x_i < l_i \\ u_i & \text{if } x_i > u_i \\ x_i & \text{otherwise.} \end{cases} \quad (4)$$

Note that $P(x_k - g_k, l, u) - x_k$ denotes the gradient projected onto the feasible design space. This convergence criteria is computed in the function `get_optimality` in *LBFGS.m*.

2.2 Computation of the Cauchy point

For simplicity, we will denote the Hessian approximation B_k , the objective gradient g_k , and the design vector x_k at the k^{th} iteration as B , g , and x , respectively. The Cauchy point x^c is parameterized as $x^c = x(t^*)$, where t^* is the first local minimizer along the piece-wise linear path $P(t) = P(x - tg, l, u)$, where P was previously defined in equation (4). Each coordinate of the piece-wise linear path $x_i(t)$ is defined as $x_i - tg_i$ $t_i \in [0, t_i]$, where the breakpoint t_i is given as:

$$t_i = \begin{cases} (x_i - u_i)/g_i & \text{if } g_i < 0, \\ (x_i - l_i)/g_i & \text{if } g_i > 0, \\ \infty & \text{otherwise.} \end{cases} \quad (5)$$

The breakpoints t_i are sorted into an ordered set $\{t^j : t^{j-1} < t^j, j = 1, 2, \dots, n\}$ and each interval $[t^{j-1}, t^j]$ is subsequently analyzed until a local minimizer t^* is found. The details of this minimization process can be found in Algorithm CP of [1], and are implemented in the routine `get_cauchy_point` in *LBFGS.m*.

2.3 Subspace minimization

Once the Cauchy point x^c has been computed, the quadratic model m_k is minimized for the free variables of x^c , those whose values are not equal to upper or lower bound values, in an unconstrained manner to obtain a solution vector d^u . This is done by directly inverting the L-BFGS-B Hessian approximation B_k using the Sherman-Morrison-Woodbury formula. Once the unconstrained minimizer is found, it is backtracked towards the feasible region with a positive scaling parameter α^* to provide a backtracked solution vector $d_k = \alpha^* d^u$. This provides enough information to compute \bar{x}_k , the minimization to the subspace problem with the bounds imposed on the design variables. The details of the subspace minimization are given in Section 5.1 of [1] and the method is implemented in *LBFGS.m* in the routine [subspace_min](#).

2.4 Strong-Wolfe line search

The new search direction is then given as $d_k = \bar{x}_k - x_k$. A line search algorithm is performed in this direction to find a step length α_k that satisfies the strong-Wolfe conditions given by the sufficient decrease condition:

$$f(x_{k+1}) \leq f(x_k) + c_1 \alpha_k g_k^T d_k, \quad (6)$$

and the curvature condition:

$$|g_{k+1}^T d_k| \leq c_2 |g_k^T d_k|, \quad (7)$$

where c_1 and c_2 are chosen to be 10^{-4} and 0.9, respectively. The algorithm to compute this line search is given by Algorithms 3.5 and 3.6 of the reference [2], and is implemented in *LBFGS.m* in the routine [strong_wolfe](#).

2.5 L-BFGS updates

Limited memory BFGS Hessian approximations store m iteration pairs $Y_k = [y_{k-m}, y_{k-m+1}, \dots, y_{k-1}]$ and $S_k = [s_{k-m}, s_{k-m+1}, \dots, s_{k-1}]$ where $y_k = g_k - g_{k-1}$ and $s_k = x_k - x_{k-1}$, where the Hessian approximation B_k can be written as:

$$B_k = \theta I - W_k M_k^{-1} W_k^T. \quad (8)$$

The expressions for the matrices W_k and M_k are outlined in Section 3 of reference [1]. The updates to the L-BFGS data structures occur inside the quasi-Newton iterations in *LBFGS.m* starting on [line 66](#).

3 Results

3.1 Convex quadratic functions

The `LBFGS.m` implementation is bundled with a series of basic [regression tests](#) that ensure the algorithm performs as expected for the two simple objective functions:

$$f_1(x) = x^T x, \tag{9}$$

and

$$f_2(x) = -x^T x. \tag{10}$$

with $n = 100$ design variables, $m = 10$ stored iteration pairs for the LBFGS data structures, and a convergence tolerance of 10^{-5} . Using these two objective functions, six regression tests are defined.

The [first regression test](#) is defined as the problem

$$\begin{aligned} & \underset{x}{\text{minimize}} && f_1(x) \\ & \text{subject to} && -10 \leq x \leq 10, \end{aligned} \tag{11}$$

with the initial starting guess $x_i = 5 \forall i$. The solution behaves essentially as an unconstrained optimization problem and converges in a single quasi-Newton iteration to the correct solution $x_i = 0 \forall i$.

The [second regression test](#) is identical to the first, with the exception that the lower bound is set to $l = 1$ rather than -10 . Again this regression test converges in a single quasi-Newton iteration to the correct solution $x_i = 1 \forall i$.

The [third regression test](#) is identical to the first, with the exception that the initial guess is defined outside of the box-constraints, such that $x_i = -20 \forall i$. This test converges in two quasi-Newton iterations to the correct solution $x_i = 0 \forall i$.

The [fourth regression test](#) is defined as

$$\begin{aligned} & \underset{x}{\text{minimize}} && f_1(x) \\ & \text{subject to} && 1 \leq x \leq 10, \end{aligned} \tag{12}$$

with the initial guess defined as a random vector whose components satisfy $9 \leq x_i \leq 10 \forall i$. This test converges to the correct solution $x_i = 1 \forall i$ in one quasi-Newton iteration.

To verify upper bounds can be properly satisfied, the [fifth regression test](#) is defined as:

$$\begin{aligned} & \underset{x}{\text{minimize}} && f_2(x) \\ & \text{subject to} && 0 \leq x \leq 10, \end{aligned} \tag{13}$$

where the initial design vector is chosen as $x_i = 5 \forall i$. This test converges correctly to the solution $x_i = 10 \forall i$ in one quasi-Newton iteration.

Finally, the [sixth test](#) is defined such that a non-uniform lower bound constraint is set, where the optimization problem is defined as:

$$\begin{aligned} & \underset{x}{\text{minimize}} && f_1(x) \\ & \text{subject to} && \sin(\pi \frac{1-i}{n}) \leq x_i \leq 10 \quad \forall i \end{aligned} \tag{14}$$

which correctly converges to the solution $x_i = \sin(\pi \frac{1-i}{n}) \forall i$ in one quasi-Newton iteration.

3.2 A bounded Rosenbrock example

As a more extensive test of the L-BFGS-B implementation, a bounded Rosenbrock [example](#) is included in the source distribution. We consider the Rosenbrock function defined as:

$$f_3(x) = 100(x_2 - x_1^2)^2 + (1 - x_1)^2, \tag{15}$$

where $x \in \mathbb{R}^2$. We first use the L-BFGS-B algorithm to solve the unconstrained minimization problem

$$\begin{aligned} & \underset{x}{\text{minimize}} && f_3(x) \\ & \text{subject to} && -\infty < x < \infty \end{aligned} \tag{16}$$

with the initial guess $(x_1, x_2) = (-1.2, 1)$, using $m = 10$ stored L-BFGS iteration pairs and a maximum number of quasi-Newton iterations of 100. The search path for the unconstrained problem is shown in [Figure 1](#). Without active bound variables, the L-BFGS-B method essentially reduces to the standard L-BFGS method. This result gives us confidence that the L-BFGS data structures and updates are performed correctly in the *LBFGS.m* implementation. The convergence history for the unconstrained Rosenbrock problem is shown in [Figure 2](#).




Figure 1: Search history for the unconstrained Rosenbrock example

As a second example, the Rosenbrock problem is modified to include a box constraint, such that the new problem definition is given as:

$$\begin{aligned}
 & \underset{x}{\text{minimize}} && f_3(x) \\
 & \text{subject to} && -0.5 \leq x_1 \leq 0.5, \\
 & && -0.5 \leq x_2 \leq 0.5,
 \end{aligned} \tag{17}$$

where the initial guess is again given as $(x_1, x_2) = (-1.2, 1.0)$. The search path produced by the *LBFGS.m* implementation is shown in Figure 3. Notice that the initial design is outside of the box constraints. The *LBFGS.m* projects the initial design guess onto the feasible design space for the first iteration, and then proceeds as normally for subsequent iterations. Notice that once the design vector enters the ‘banana valley’, it takes approximately the same path towards the global minimum that the unbounded Rosenbrock example takes. However, it finds a bounded minimum at the design location $x = (0.5, 0.25)$. This result was corroborated by comparison to a result obtained via Matlab’s robust optimization solver *fmincon*. The convergence history for the bounded Rosenbrock example is shown in Figure 4.

iter	f(x)	optimality
0	24.20000000	215.60000000
1	5.10111266	38.33803031
2	4.15378843	6.65579078
3	4.11721504	1.40757360
4	4.10921995	1.42455138
5	4.10922456	1.42457526
6	3.58186385	11.18575152
7	3.45470246	17.47839459
8	3.28544194	16.76380586
9	2.65947615	10.87782125
10	2.22970801	5.83940028
11	2.09539857	8.36664503
12	1.75072527	11.07438171
13	1.39706792	2.79021511
14	1.11088550	4.82357799
15	1.03878169	8.04442772
16	0.79586966	3.26617858
17	0.64945895	6.65883593
18	0.46730756	1.11095742
19	0.38120958	3.61380605
20	0.30770768	5.10591921
21	0.19454187	0.57434669
22	0.15033422	2.04613244
23	0.12796165	6.48492562
24	0.07920583	0.77105145
25	0.05289519	0.49482493
26	0.03442939	2.68276032
27	0.01970635	1.93274402
28	0.00863828	0.32795662
29	0.00470913	1.91555193
30	0.00202819	0.48371906
31	0.00035537	0.22332188
32	0.00004326	0.18103267
33	0.00000213	0.03578927
34	0.00000012	0.00998035
35	0.00000000	0.00006825
36	0.00000000	0.00000511

Figure 2: Convergence history for the unbounded Rosenbrock example

4 Conclusions

We have implemented a pure Matlab implementation of the L-BFGS-B algorithm for box-constrained gradient-based optimization. We have provided a high-level overview of the fundamental steps of the algorithm and provided references to the corresponding implementation locations of these steps in the *LBFGS.m* code. The implementation has been verified for a series of simple tests. Further work includes verifying the implementation for a wider variety of optimization problems.

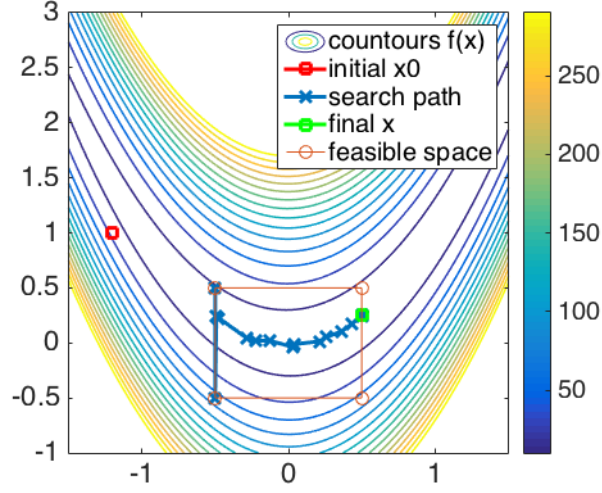


Figure 3: Search history for constrained Rosenbrock example

iter	f(x)	optimality
0	8.50000000	1.00000000
1	58.50000000	1.00000000
2	2.22755657	0.99250000
3	2.17219686	0.97383645
4	1.78942574	0.78413677
5	1.58869314	0.72538530
6	1.26415453	0.62376555
7	1.08833255	0.53790684
8	0.93554003	0.51145792
9	0.71623471	0.71067169
10	0.56619271	0.44590189
11	0.47950989	0.85958247
12	0.34314773	0.93149460
13	0.25780652	0.76709009
14	0.25011825	0.21748547
15	0.25000000	0.00000000

Figure 4: Convergence history for the unbounded Rosenbrock example

References

- [1] Byrd, R. H., Lu, P., Nocedal, J., & Zhu, C. (1995). *A limited memory algorithm for bound constrained optimization*. SIAM Journal on Scientific Computing, 16(5), 1190-1208.
- [2] Nocedal, J. and Wright, S. *Numerical optimization*. Springer Science &

Business Media, 2006.

- [3] Byrd, R. H., Lu, P., Nocedal, J., & Zhu, C. *L-BFGS-B*. <http://users.iems.northwestern.edu/~nocedal/lbfgsb.html>. Accessed: 2016-12-13.
- [4] Becker, S. *L-BFGS-B-C* <https://github.com/stephenbecker/L-BFGS-B-C>. Accessed: 2016-12-13.
- [5] Carbonetto, P. *lbfgsb-matlab*. <https://github.com/pcarbo/lbfgsb-matlab>. Accessed: 2016-12-13.
- [6] Kobos, M. *lbfgsb-wrapper* <https://github.com/mkobos/lbfgsb-wrapper>. Accessed: 2016-12-13.
- [7] Yu, H. *Lbfgsb.jl*. <https://github.com/yuhonglin/Lbfgsb.jl>. Accessed: 2016-12-13.
- [8] Byrd, Richard H., Jorge Nocedal, and Robert B. Schnabel. *Representations of quasi-Newton matrices and their use in limited memory methods*. Mathematical Programming 63.1-3 (1994): 129-156.